References

1. Jinek, M., Chylinski, K., Fonfara, I., Hauer, M., Doudna, J. A. & Charpentier, E. A programmable dual-RNA–guided DNA endonuclease in adaptive bacterial immunity. *science* **337,** 816–821 (2012).
2. Hopfield, J. J. Neural networks and physical systems with emergent collective computational abilities. *Proceedings of the national academy of sciences* **79,** 2554–2558 (1982).
3. Hinton, G. E., Sejnowski, T. J., *et al.* Learning and relearning in Boltzmann machines. *Parallel distributed processing: Explorations in the microstructure of cognition* **1,** 2 (1986).
4. Guo, D., Yang, D., Zhang, H., Song, J., Zhang, R., Xu, R., Zhu, Q., Ma, S., Wang, P., Bi, X., *et al.* Deepseek-r1: Incentivizing reasoning capability in llms via reinforcement learning. *arXiv preprint arXiv:2501.12948* (2025).
5. Jaech, A., Kalai, A., Lerer, A., Richardson, A., El-Kishky, A., Low, A., Helyar, A., Madry, A., Beutel, A., Carney, A., *et al.* Openai o1 system card. *arXiv preprint arXiv:2412.16720* (2024).
6. Team, G., Georgiev, P., Lei, V. I., Burnell, R., Bai, L., Gulati, A., Tanzer, G., Vincent, D., Pan, Z., Wang, S., *et al.* Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context. *arXiv preprint arXiv:2403.05530* (2024).
7. Wiesinger, J., Marlow, P. & Vuskovic, V. Agents (2024).
8. Hinton, G. Distilling the Knowledge in a Neural Network. *arXiv preprint arXiv:1503.02531* (2015).
9. Team, G., Mesnard, T., Hardin, C., Dadashi, R., Bhupatiraju, S., Pathak, S., Sifre, L., Riviére, M., Kale, M. S., Love, J., *et al.* Gemma: Open models based on gemini research and technology. *arXiv preprint arXiv:2403.08295* (2024).
10. Leslie, D., Ashurst, C., González, N. M., Griffiths, F., Jayadeva, S., Jorgensen, M., Katell, M., Krishna, S., Kwiatkowski, D., Martins, C. I., *et al.* 'Frontier AI,'Power, and the Public Interest: Who benefits, who decides? *Harvard Data Science Review* (2024).
11. Chen, L., Davis, J. Q., Hanin, B., Bailis, P., Stoica, I., Zaharia, M. & Zou, J. *Are More LLM Calls All You Need? Towards the Scaling Properties of Compound AI Systems* in *The Thirty-eighth Annual Conference on Neural Information Processing Systems* (2024).
12. Gower, B. *Scientific method: A historical and philosophical introduction* (Routledge, 2012).
13. Snell, C., Lee, J., Xu, K. & Kumar, A. Scaling llm test-time compute optimally can be more effective than scaling model parameters. *arXiv preprint arXiv:2408.03314* (2024).
14. Brown, N. & Sandholm, T. Superhuman AI for multiplayer poker. *Science* **365,** 885–890 (2019).
15. Silver, D., Huang, A., Maddison, C. J., Guez, A., Sifre, L., Van Den Driessche, G., Schrittwieser, J., Antonoglou, I., Panneershelvam, V., Lanctot, M., *et al.* Mastering the game of Go with deep neural networks and tree search. *Nature* **529,** 484–489 (2016).
16. Ringel, M. S., Scannell, J. W., Baedeker, M. & Schulze, U. Breaking Eroom's law. *Nat Rev Drug Discov* **19,** 833–834 (2020).
17. Pushpakom, S., Iorio, F., Eyers, P. A., Escott, K. J., Hopper, S., Wells, A., Doig, A., Guilliams, T., Latimer, J., McNamee, C., *et al.* Drug repurposing: progress, challenges and recommendations. *Nature reviews Drug discovery* **18,** 41–58 (2019).
18. Xia, Y., Sun, M., Huang, H. & Jin, W.-L. Drug repurposing for cancer therapy. *Signal Transduction and Targeted Therapy* **9,** 92. https://doi.org/10.1038/s41392-024-01808-1 (Apr. 2024).
19. Keown, O. P., Warburton, W., Davies, S. C. & Darzi, A. Antimicrobial resistance: addressing the global threat through greater awareness and transformative action. *Health Affairs* **33,** 1620–1626 (2014).
20. He, L., Patkowski, J. B., Miguel-Romero, L., Aylett, C. H., Fillol-Salom, A., Costa, T. R. & Penades, J. R. Chimeric infective particles expand species boundaries in phage inducible chromosomal island mobilization. *bioRxiv,* 2025–02 (2025).
21. Penadés, J. R., Gottweis, J., He, L., Patkowski, J. B., Shurick, A., Weng, W.-H., Tu, T., Palepu, A., Myaskovsky, A., Pawlosky, A., Natarajan, V., Karthikesalingam, A. & Costa, T. R. D. *AI mirrors experimental science to uncover a novel mechanism of gene transfer crucial to bacterial evolution* https://storage.googleapis.com/coscientist_paper/penades2025ai.pdf.
22. Bommasani, R., Hudson, D. A., Adeli, E., Altman, R., Arora, S., von Arx, S., Bernstein, M. S., Bohg, J., Bosselut, A., Brunskill, E., *et al.* On the opportunities and risks of foundation models. *arXiv preprint arXiv:2108.07258* (2021).
23. Erhan, D., Courville, A., Bengio, Y. & Vincent, P. *Why does unsupervised pre-training help deep learning?* in *Proceedings of the thirteenth international conference on artificial intelligence and statistics* (2010), 201–208.
24. Radford, A. Improving language understanding by generative pre-training (2018).
25. Team, G., Anil, R., Borgeaud, S., Wu, Y., Alayrac, J.-B., Yu, J., Soricut, R., Schalkwyk, J., Dai, A. M., Hauth, A., *et al.* Gemini: a family of highly capable multimodal models. *arXiv preprint arXiv:2312.11805* (2023).
26. Achiam, J., Adler, S., Agarwal, S., Ahmad, L., Akkaya, I., Aleman, F. L., Almeida, D., Altenschmidt, J., Altman, S., Anadkat, S., *et al.* GPT-4 technical report. *arXiv preprint arXiv:2303.08774* (2023).
27. Chowdhery, A., Narang, S., Devlin, J., Bosma, M., Mishra, G., Roberts, A., Barham, P., Chung, H. W., Sutton, C., Gehrmann, S., *et al.* PaLM: Scaling language modeling with pathways. *arXiv preprint arXiv:2204.02311* (2022).
28. Google. *PaLM 2 Technical Report* https://ai.google/static/documents/palm2techreport.pdf. 2023.
29. Wei, J., Wang, X., Schuurmans, D., Bosma, M., Xia, F., Chi, E., Le, Q. V., Zhou, D., *et al.* Chain-of-thought prompting elicits reasoning in large language models. *Advances in neural information processing systems* **35,** 24824–24837 (2022).
30. Kahneman, D. Thinking, fast and slow. *Farrar, Straus and Giroux* (2011).
31. Yao, S., Yu, D., Zhao, J., Shafran, I., Griffiths, T., Cao, Y. & Narasimhan, K. Tree of thoughts: Deliberate problem solving with large language models. *Advances in Neural Information Processing Systems* **36** (2024).
32. Zelikman, E., Wu, Y., Mu, J. & Goodman, N. Star: Bootstrapping reasoning with reasoning. *Advances in Neural Information Processing Systems* **35,** 15476–15488 (2022).
33. Williams, T., Szekendi, M., Pavkovic, S., Clevenger, W. & Cerese, J. The reliability of AHRQ Common Format Harm Scales in rating patient safety events. eng. *J Patient Saf* **11,** 52–59. ISSN: 1549-8425 (Mar. 2015).
34. Muennighoff, N., Yang, Z., Shi, W., Li, X. L., Fei-Fei, L., Hajishirzi, H., Zettlemoyer, L., Liang, P., Candés, E. & Hashimoto, T. s1: Simple test-time scaling. *arXiv preprint arXiv:2501.19393* (2025).

35. Tu, T., Palepu, A., Schaekermann, M., Saab, K., Freyberg, J., Tanno, R., Wang, A., Li, B., Amin, M., Tomasev, N., *et al.* Towards conversational diagnostic AI. *arXiv preprint arXiv:2401.05654* (2024).

36. Jumper, J., Evans, R., Pritzel, A., Green, T., Figurnov, M., Ronneberger, O., Tunyasuvunakool, K., Bates, R., Žídek, A., Potapenko, A., *et al.* Highly accurate protein structure prediction with AlphaFold. *Nature* **596,** 583–589 (2021).

37. Wong, F., Zheng, E. J., Valeri, J. A., Donghia, N. M., Anahtar, M. N., Omori, S., Li, A., Cubillos-Ruiz, A., Krishnan, A., Jin, W., *et al.* Discovery of a structural class of antibiotics with explainable deep learning. *Nature* **626,** 177–185 (2024).

38. Zambaldi, V., La, D., Chu, A. E., Patani, H., Danson, A. E., Kwan, T. O., Frerix, T., Schneider, R. G., Saxton, D., Thillaisundaram, A., *et al.* De novo design of high-affinity protein binders with AlphaProteo. *arXiv preprint arXiv:2409.08022* (2024).

39. Merchant, A., Batzner, S., Schoenholz, S. S., Aykol, M., Cheon, G. & Cubuk, E. D. Scaling deep learning for materials discovery. *Nature* **624,** 80–85 (2023).

40. Lu, C., Lu, C., Lange, R. T., Foerster, J., Clune, J. & Ha, D. The AI scientist: Towards fully automated open-ended scientific discovery. *arXiv preprint arXiv:2408.06292* (2024).

41. Schmidgall, S., Su, Y., Wang, Z., Sun, X., Wu, J., Yu, X., Liu, J., Liu, Z. & Barsoum, E. Agent laboratory: Using llm agents as research assistants. *arXiv preprint arXiv:2501.04227* (2025).

42. Liang, W., Zhang, Y., Cao, H., Wang, B., Ding, D. Y., Yang, X., Vodrahalli, K., He, S., Smith, D. S., Yin, Y., *et al.* Can large language models provide useful feedback on research papers? A large-scale empirical analysis. *NEJM AI* **1,** AIoa2400196 (2024).

43. OpenAI. *GPT-4 Technical Report* 2023. arXiv: 2303.08774 [cs.CL].

44. Skarlinski, M. D., Cox, S., Laurent, J. M., Braza, J. D., Hinks, M., Hammerling, M. J., Ponnapati, M., Rodriques, S. G. & White, A. D. Language agents achieve superhuman synthesis of scientific knowledge. *arXiv preprint arXiv:2409.13740* (2024).

45. Zhou, D., Schärli, N., Hou, L., Wei, J., Scales, N., Wang, X., Schuurmans, D., Bousquet, O., Le, Q. & Chi, E. Least-to-Most Prompting Enables Complex Reasoning in Large Language Models. *arXiv preprint arXiv:2205.10625* (2022).

46. Ifargan, T., Hafner, L., Kern, M., Alcalay, O. & Kishony, R. Autonomous LLM-Driven Research—from Data to Human-Verifiable Research Papers. *NEJM AI* **2,** AIoa2400555 (2025).

47. Swanson, K., Wu, W., Bulaong, N. L., Pak, J. E. & Zou, J. The virtual lab: AI agents design new SARS-CoV-2 nanobodies with experimental validation. *bioRxiv,* 2024–11 (2024).

48. Boiko, D. A., MacKnight, R., Kline, B. & Gomes, G. Autonomous chemical research with large language models. *Nature* **624,** 570–578 (2023).

49. Singhal, K., Azizi, S., Tu, T., Mahdavi, S. S., Wei, J., Chung, H. W., Scales, N., Tanwani, A., Cole-Lewis, H., Pfohl, S., *et al.* Large Language Models Encode Clinical Knowledge. *arXiv preprint arXiv:2212.13138* (2022).

50. Saab, K., Tu, T., Weng, W.-H., Tanno, R., Stutz, D., Wulczyn, E., Zhang, F., Strother, T., Park, C., Vedadi, E., *et al.* Capabilities of gemini models in medicine. *arXiv preprint arXiv:2404.18416* (2024).

51. Taylor, R., Kardas, M., Cucurull, G., Scialom, T., Hartshorn, A., Saravia, E., Poulton, A., Kerkez, V. & Stojnic, R. Galactica: A Large Language Model for Science. *arXiv preprint arXiv:2211.09085* (2022).

52. Chaves, J. M. Z., Wang, E., Tu, T., Vaishnav, E. D., Lee, B., Mahdavi, S. S., Semturs, C., Fleet, D., Natarajan, V. & Azizi, S. Tx-LLM: A Large Language Model for Therapeutics. *arXiv preprint arXiv:2406.06316* (2024).

53. Tu, T., Fang, Z., Cheng, Z., Spasic, S., Palepu, A., Stankovic, K. M., Natarajan, V. & Peltz, G. Genetic Discovery Enabled by A Large Language Model. *bioRxiv* (2023).

54. Nguyen, E., Poli, M., Durrant, M. G., Kang, B., Katrekar, D., Li, D. B., Bartie, L. J., Thomas, A. W., King, S. H., Brixi, G., *et al.* Sequence modeling and design from molecular to genome scale with Evo. *Science* **386,** eado9336 (2024).

55. Lin, Z., Akin, H., Rao, R., Hie, B., Zhu, Z., Lu, W., Smetanin, N., Verkuil, R., Kabeli, O., Shmueli, Y., *et al.* Evolutionary-scale prediction of atomic-level protein structure with a language model. *Science* **379,** 1123–1130 (2023).

56. Ruffolo, J. A., Nayfach, S., Gallagher, J., Bhatnagar, A., Beazer, J., Hussain, R., Russ, J., Yip, J., Hill, E., Pacesa, M., *et al.* Design of highly functional genome editors by modeling the universe of CRISPR-Cas sequences. *bioRxiv,* 2024–04 (2024).

57. Shaw, P., Gurram, B., Belanger, D., Gane, A., Bileschi, M. L., Colwell, L. J., Toutanova, K. & Parikh, A. P. ProtEx: A Retrieval-Augmented Approach for Protein Function Prediction. *bioRxiv,* 2024–05 (2024).

58. Krishnamurthy, N., Grimshaw, A. A., Axson, S. A., Choe, S. H. & Miller, J. E. Drug repurposing: a systematic review on root causes, barriers and facilitators. *BMC health services research* **22,** 970 (2022).

59. Zitnik, M., Agrawal, M. & Leskovec, J. Modeling polypharmacy side effects with graph convolutional networks. *Bioinformatics* **34,** i457–i466 (2018).

60. Morselli Gysi, D., Do Valle, Í., Zitnik, M., Ameli, A., Gan, X., Varol, O., Ghiassian, S. D., Patten, J., Davey, R. A., Loscalzo, J., *et al.* Network medicine framework for identifying drug-repurposing opportunities for COVID-19. *Proceedings of the National Academy of Sciences* **118,** e2025581118 (2021).

61. Huang, K., Chandak, P., Wang, Q., Havaldar, S., Vaid, A., Leskovec, J., Nadkarni, G., Glicksberg, B. S., Gehlenborg, N. & Zitnik, M. A foundation model for clinician-centered drug repurposing. *medRxiv* (2024).

62. Jones, N. OpenAI's 'deep research' tool: is it useful for scientists? *Nature News.* https://doi.org/10.1038/d41586-025-00377-9 (2025).

63. Megat, S., Mora, N., Sanogo, J., Roman, O., Catanese, A., Alami, N. O., Freischmidt, A., Mingaj, X., De Calbiac, H., Muratet, F., *et al.* Integrative genetic analysis illuminates ALS heritability and identifies risk genes. *Nature communications* **14,** 342 (2023).

64. Elo, A. E. & Sloan, S. The rating of chessplayers: Past and present. *(No Title)* (1978).

65. LeCun, Y., Touresky, D., Hinton, G. & Sejnowski, T. *A theoretical framework for back-propagation* in *Proceedings of the 1988 connectionist models summer school* **1** (1988), 21–28.

66. Hamilton, A. H., Roughan, M. & Kalenkova, A. Elo Ratings in the Presence of Intransitivity. *arXiv preprint arXiv:2412.14427* (2024).

67. Coulom, R. Computing "elo ratings" of move patterns in the game of go. *ICGA journal* **30,** 198–208 (2007).

68. Kovalchik, S. Extension of the Elo rating system to margin of victory. *International Journal of Forecasting* **36,** 1329–1341 (2020).

69. Khan, A., Hughes, J., Valentine, D., Ruis, L., Sachan, K., Radhakrishnan, A., Grefenstette, E., Bowman, S. R., Rocktäschel, T. & Perez, E. Debating with more persuasive llms leads to more truthful answers. *arXiv preprint arXiv:2402.06782* (2024).
70. Post, M. & Vilar, D. Fast lexically constrained decoding with dynamic beam allocation for neural machine translation. *arXiv preprint arXiv:1804.06609* (2018).
71. Rein, D., Hou, B. L., Stickland, A. C., Petty, J., Pang, R. Y., Dirani, J., Michael, J. & Bowman, S. R. GPQA: A graduate-level google-proof q&a benchmark. *arXiv preprint arXiv:2311.12022* (2023).
72. *DepMap Q2 2024 Data Release* 2024. https://depmap.org/portal/.
73. Döhner, H., Weisdorf, D. J. & Bloomfield, C. D. Acute myeloid leukemia. *New England Journal of Medicine* **373,** 1136–1152 (2015).
74. Guo, Q., Jin, Y., Chen, X., Ye, X., Shen, X., Lin, M., Zeng, C., Zhou, T. & Zhang, J. NF-κB in biology and targeted therapy: new insights and translational implications. *Signal Transduction and Targeted Therapy* **9,** 53 (2024).
75. Van Herpen, C. M., Agarwala, S. S., Hauschild, A., Berking, C., Beck, J. T., Schadendorf, D., Jansen, R., Queirolo, P., Ascierto, P. A., Blank, C. U., *et al.* Biomarker results from a phase II study of MEK1/2 inhibitor binimetinib (MEK162) in patients with advanced NRAS-or BRAF-mutated melanoma. *Oncotarget* **10,** 1850 (2019).
76. Ball, B. J., Hsu, M., Devlin, S. M., Arcila, M., Roshal, M., Zhang, Y., Famulare, C. A., Goldberg, A. D., Cai, S. F., Dunbar, A., *et al.* The prognosis and durable clearance of RAS mutations in patients with acute myeloid leukemia receiving induction chemotherapy. *American journal of hematology* **96,** E171 (2021).
77. Verstovsek, S. & Komrokji, R. S. A comprehensive review of pacritinib in myelofibrosis. *Future Oncology* **11,** 2819–2830 (2015).
78. Perrone, S., Ottone, T., Zhdanovskaya, N. & Molica, M. How acute myeloid leukemia (AML) escapes from FMS-related tyrosine kinase 3 (FLT3) inhibitors? Still an overrated complication? *Cancer Drug Resistance* **6,** 223 (2023).
79. Bosman, M. C. J., Schuringa, J. J. & Vellenga, E. Constitutive NF-κB activation in AML: Causes and treatment strategies. *Critical reviews in oncology/hematology* **98,** 35–44 (2016).
80. Hou, S., Wang, X., Guo, T., Lan, Y., Yuan, S., Yang, S., Zhao, F., Fang, A., Liu, N., Yang, W., *et al.* PHF6 maintains acute myeloid leukemia via regulating NF-κB signaling pathway. *Leukemia* **37,** 1626–1637 (2023).
81. Banker, D. E., Mayer, S. J., Li, H. Y., Willman, C. L., Appelbaum, F. R. & Zager, R. A. Cholesterol synthesis and import contribute to protective cholesterol increments in acute myeloid leukemia cells. *Blood* **104,** 1816–1824 (2004).
82. Guan, Y., Enejder, A., Wang, M., Fang, Z., Cui, L., Chen, S.-Y., Wang, J., Tan, Y., Wu, M., Chen, X., *et al.* A human multi-lineage hepatic organoid model for liver fibrosis. *Nature communications* **12,** 6138 (2021).
83. Guan, Y., Fang, Z., Hu, A., Roberts, S., Wang, M., Ren, W., Johansson, P. K., Heilshorn, S. C., Enejder, A. & Peltz, G. Live Cell Imaging of Human Liver Fibrosis using Hepatic Micro-Organoids. *JCI insight* (2024).
84. Guan, Y. & Peltz, G. Hepatic organoids move from adolescence to maturity. *Liver International* **44,** 1290–1297 (2024).
85. He, J. & Vechev, M. *Large language models for code: Security hardening and adversarial testing* in *Proceedings of the 2023 ACM SIGSAC Conference on Computer and Communications Security* (2023), 1865–1879.
86. Alqurainy, N., Miguel-Romero, L., de Sousa, J. M., Chen, J., Rocha, E. P., Fillol-Salom, A. & Penadés, J. R. A widespread family of phage-inducible chromosomal islands only steals bacteriophage tails to spread in nature. *Cell Host & Microbe* **31,** 69–82 (2023).
87. De Sousa, J. A. M., Fillol-Salom, A., Penadés, J. R. & Rocha, E. P. Identification and characterization of thousands of bacteriophage satellites across bacteria. *Nucleic acids research* **51,** 2759–2777 (2023).
88. Brazil, R. Illuminating 'the ugly side of science': fresh incentives for reporting negative results. *Nature* (2024).
89. Roberts, J., Han, K., Houlsby, N. & Albanie, S. SciFIBench: Benchmarking Large Multimodal Models for Scientific Figure Interpretation. *arXiv preprint arXiv:2405.08807* (2024).
90. Shrader-Frechette, K. S. *Ethics of scientific research* (Rowman & Littlefield, 1994).
91. Resnik, D. B. *The ethics of science: An introduction* (Routledge, 2005).
92. Rollin, B. E. *Science and ethics* (Cambridge University Press, 2006).
93. Fisher, C. B. & Anushko, A. E. Research ethics in social science. *The SAGE handbook of social research methods,* 95–109 (2008).
94. Edel, A. *Science and the structure of ethics* (Routledge, 2018).
95. Menapace, M. Scientific ethics: A new approach. *Science and Engineering Ethics* **25,** 1193–1216 (2019).
96. Miller, S. & Selgelid, M. J. Ethical and philosophical consideration of the dual-use dilemma in the biological sciences. *Science and engineering ethics* **13,** 523–580 (2007).
97. Selgelid, M. J. Governance of dual-use research: an ethical dilemma. *Bulletin of the World Health Organization* **87,** 720–723 (2009).
98. Pustovit, S. V. & Williams, E. D. Philosophical aspects of dual use technologies. *Science and engineering ethics* **16,** 17–31 (2010).
99. Forge, J. A note on the definition of "dual use". *Science and Engineering Ethics* **16,** 111–118 (2010).
100. Kuhlau, F., Höglund, A. T., Eriksson, S. & Evers, K. The ethics of disseminating dual-use knowledge. *Research Ethics* **9,** 6–19 (2013).
101. Shaw, S. & Barrett, G. Research governance: regulating risk and reducing harm? *Journal of the Royal Society of Medicine* **99,** 14–19 (2006).
102. Rothstein, H., Irving, P., Walden, T. & Yearsley, R. The risks of risk-based regulation: Insights from the environmental policy domain. *Environment international* **32,** 1056–1065 (2006).
103. Ludlow, K., Bowman, D. M., Gatof, J. & Bennett, M. G. Regulating emerging and future technologies in the present. *NanoEthics* **9,** 151–163 (2015).
104. Verschraegen, G. Regulating scientific research: A constitutional moment? *Journal of Law and Society* **45,** S163–S184 (2018).
105. Jobin, A., Ienca, M. & Vayena, E. The global landscape of AI ethics guidelines. *Nature machine intelligence* **1,** 389–399 (2019).

106. Wansley, M. T. Regulation of Emerging Risks. *Vand. L. Rev.* **69,** 401 (2016).

107. Gabriel, I., Manzini, A., Keeling, G., Hendricks, L. A., Rieser, V., Iqbal, H., Tomašev, N., Ktena, I., Kenton, Z., Rodriguez, M., *et al.* The ethics of advanced ai assistants. *arXiv preprint arXiv:2404.16244* (2024).

108. Tang, X., Jin, Q., Zhu, K., Yuan, T., Zhang, Y., Zhou, W., Qu, M., Zhao, Y., Tang, J., Zhang, Z., Cohan, A., Lu, Z. & Gerstein, M. *Prioritizing Safeguarding Over Autonomy: Risks of LLM Agents for Science* 2024. arXiv: 2402.04247 [cs.CY]. https://arxiv.org/abs/2402.04247.

109. Shevlane, T., Farquhar, S., Garfinkel, B., Phuong, M., Whittlestone, J., Leung, J., Kokotajlo, D., Marchal, N., Anderljung, M., Kolt, N., *et al.* Model evaluation for extreme risks. *arXiv preprint arXiv:2305.15324* (2023).

110. Bova, P., Di Stefano, A. & Han, T. A. Quantifying detection rates for dangerous capabilities: a theoretical model of dangerous capability evaluations. *arXiv preprint arXiv:2412.15433* (2024).

111. Phuong, M., Aitchison, M., Catt, E., Cogan, S., Kaskasoli, A., Krakovna, V., Lindner, D., Rahtz, M., Assael, Y., Hodkinson, S., *et al.* Evaluating frontier models for dangerous capabilities. *arXiv preprint arXiv:2403.13793* (2024).

112. Sabour, S., Liu, J. M., Liu, S., Yao, C. Z., Cui, S., Zhang, X., Zhang, W., Cao, Y., Bhat, A., Guan, J., Wu, W., Mihalcea, R., Althoff, T., Lee, T. M. C. & Huang, M. *Human Decision-making is Susceptible to AI-driven Manipulation* 2025. arXiv: 2502.07663 [cs.AI]. https://arxiv.org/abs/2502.07663.

113. Ke, P., Wen, B., Feng, Z., Liu, X., Lei, X., Cheng, J., Wang, S., Zeng, A., Dong, Y., Wang, H., *et al.* Critiquellm: Scaling llm-as-critic for effective and explainable evaluation of large language model generation. *arXiv preprint arXiv:2311.18702* (2023).

114. Vu, T., Krishna, K., Alzubi, S., Tar, C., Faruqui, M. & Sung, Y.-H. Foundational autoraters: Taming large language models for better automatic evaluation. *arXiv preprint arXiv:2407.10817* (2024).

115. Wei, H., He, S., Xia, T., Wong, A., Lin, J. & Han, M. Systematic evaluation of llm-as-a-judge in llm alignment tasks: Explainable metrics and diverse prompt templates. *arXiv preprint arXiv:2408.13006* (2024).

116. Lan, T., Zhang, W., Xu, C., Huang, H., Lin, D., Chen, K. & Mao, X.-l. Criticbench: Evaluating large language models as critic. *arXiv preprint arXiv:2402.13764* (2024).

117. Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E., *et al.* Judging llm-as-a-judge with mt-bench and chatbot arena. *Advances in Neural Information Processing Systems* **36** (2024).

118. Gu, J., Jiang, X., Shi, Z., Tan, H., Zhai, X., Xu, C., Li, W., Shen, Y., Ma, S., Liu, H., *et al.* A Survey on LLM-as-a-Judge. *arXiv preprint arXiv:2411.15594* (2024).

119. Gou, Z., Shao, Z., Gong, Y., Shen, Y., Yang, Y., Duan, N. & Chen, W. Critic: Large language models can self-correct with tool-interactive critiquing. *arXiv preprint arXiv:2305.11738* (2023).

120. Chen, D., Chen, R., Zhang, S., Liu, Y., Wang, Y., Zhou, H., Zhang, Q., Wan, Y., Zhou, P. & Sun, L. Mllm-as-a-judge: Assessing multimodal llm-as-a-judge with vision-language benchmark. *arXiv preprint arXiv:2402.04788* (2024).

121. Szymanski, A., Ziems, N., Eicher-Miller, H. A., Li, T. J.-J., Jiang, M. & Metoyer, R. A. *Limitations of the LLM-as-a-Judge Approach for Evaluating LLM Outputs in Expert Knowledge Tasks* 2024. arXiv: 2410.20266 [cs.HC]. https://arxiv.org/abs/2410.20266.

122. Shayegani, E., Mamun, M. A. A., Fu, Y., Zaree, P., Dong, Y. & Abu-Ghazaleh, N. Survey of vulnerabilities in large language models revealed by adversarial attacks. *arXiv preprint arXiv:2310.10844* (2023).

123. Zhu, K., Wang, J., Zhou, J., Wang, Z., Chen, H., Wang, Y., Yang, L., Ye, W., Zhang, Y., Zhenqiang Gong, N., *et al.* Promptbench: Towards evaluating the robustness of large language models on adversarial prompts. *arXiv e-prints,* arXiv–2306 (2023).

124. Fu, X., Wang, Z., Li, S., Gupta, R. K., Mireshghallah, N., Berg-Kirkpatrick, T. & Fernandes, E. Misusing tools in large language models with visual adversarial examples. *arXiv preprint arXiv:2310.03185* (2023).

125. Zhang, Z., Yang, J., Ke, P., Mi, F., Wang, H. & Huang, M. Defending large language models against jailbreaking attacks through goal prioritization. *arXiv preprint arXiv:2311.09096* (2023).

126. Chao, P., Debenedetti, E., Robey, A., Andriushchenko, M., Croce, F., Sehwag, V., Dobriban, E., Flammarion, N., Pappas, G. J., Tramer, F., *et al.* Jailbreakbench: An open robustness benchmark for jailbreaking large language models. *arXiv preprint arXiv:2404.01318* (2024).

127. Zhao, Y., Pang, T., Du, C., Yang, X., Li, C., Cheung, N.-M. M. & Lin, M. On evaluating adversarial robustness of large vision-language models. *Advances in Neural Information Processing Systems* **36** (2024).

128. Ma, X., Gao, Y., Wang, Y., Wang, R., Wang, X., Sun, Y., Ding, Y., Xu, H., Chen, Y., Zhao, Y., Huang, H., Li, Y., Zhang, J., Zheng, X., Bai, Y., Wu, Z., Qiu, X., Zhang, J., Li, Y., Sun, J., Wang, C., Gu, J., Wu, B., Chen, S., Zhang, T., Liu, Y., Gong, M., Liu, T., Pan, S., Xie, C., Pang, T., Dong, Y., Jia, R., Zhang, Y., Ma, S., Zhang, X., Gong, N., Xiao, C., Erfani, S., Li, B., Sugiyama, M., Tao, D., Bailey, J. & Jiang, Y.-G. *Safety at Scale: A Comprehensive Survey of Large Model Safety* 2025. arXiv: 2502.05206 [cs.CR]. https://arxiv.org/abs/2502.05206.

129. Zou, A., Wang, Z., Carlini, N., Nasr, M., Kolter, J. Z. & Fredrikson, M. Universal and transferable adversarial attacks on aligned language models. *arXiv preprint arXiv:2307.15043* (2023).

130. Lapid, R., Langberg, R. & Sipper, M. Open sesame! universal black box jailbreaking of large language models. *arXiv preprint arXiv:2309.01446* (2023).

131. Wang, J., Liu, Z., Park, K. H., Jiang, Z., Zheng, Z., Wu, Z., Chen, M. & Xiao, C. Adversarial demonstration attacks on large language models. *arXiv preprint arXiv:2305.14950* (2023).

132. Qiang, Y., Zhou, X. & Zhu, D. Hijacking large language models via adversarial in-context learning. *arXiv preprint arXiv:2311.09948* (2023).

133. Qi, X., Huang, K., Panda, A., Wang, M. & Mittal, P. Visual adversarial examples jailbreak large language models. *arXiv preprint arXiv:2306.13213* (2023).

134. Chao, P., Robey, A., Dobriban, E., Hassani, H., Pappas, G. J. & Wong, E. Jailbreaking black box large language models in twenty queries. *arXiv preprint arXiv:2310.08419* (2023).

135. Wang, Z., Xie, W., Wang, B., Wang, E., Gui, Z., Ma, S. & Chen, K. *Foot In The Door: Understanding Large Language Model Jailbreaking via Cognitive Psychology* 2024. arXiv: 2402.15690 [cs.CL]. https://arxiv.org/abs/2402.15690.

136. Zhang, J., Ma, X., Wang, X., Qiu, L., Wang, J., Jiang, Y.-G. & Sang, J. *Adversarial prompt tuning for vision-language models* in *European Conference on Computer Vision* (2024), 56–72.

137. Shneiderman, B. Bridging the gap between ethics and practice: guidelines for reliable, safe, and trustworthy human-centered AI systems. *ACM Transactions on Interactive Intelligent Systems (TiiS)* **10,** 1–31 (2020).

138. Anthropic. *Responsible Scaling Policy* 2024. https://assets.anthropic.com/m/24a47b00f10301cd/original/Anthropic-Responsible-Scaling-Policy-2024-10-15.pdf.

139. Anderson, B. R., Shah, J. H. & Kreminski, M. *Homogenization Effects of Large Language Models on Human Creative Ideation* in *Proceedings of the 16th Conference on Creativity & Cognition* (Association for Computing Machinery, Chicago, IL, USA, 2024), 413–425. ISBN: 9798400704857. https://doi.org/10.1145/3635636.3656204.

140. Ashkinaze, J., Mendelsohn, J., Qiwei, L., Budak, C. & Gilbert, E. *How AI Ideas Affect the Creativity, Diversity, and Evolution of Human Ideas: Evidence From a Large, Dynamic Experiment* 2024. arXiv: 2401.13481 [cs.CY]. https://arxiv.org/abs/2401.13481.

141. Wenger, E. & Kenett, Y. *We're Different, We're the Same: Creative Homogeneity Across LLMs* 2025. arXiv: 2501.19361 [cs.CY]. https://arxiv.org/abs/2501.19361.

142. Sun, N. & Amon, J. J. Addressing inequity: neglected tropical diseases and human rights. *Health and human rights* **20,** 11 (2018).

143. George, N. S., David, S. C., Nabiryo, M., Sunday, B. A., Olanrewaju, O. F., Yangaza, Y. & Shomuyiwa, D. O. Addressing neglected tropical diseases in Africa: a health equity perspective. *Global Health Research and Policy* **8,** 30 (2023).